

```
# Import required libraries
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('seaborn-v0_8')
%matplotlib inline
```

▼ Bank Customer Segmentation Analysis - Student Exercise

In this lab, you will implement customer segmentation using K-means clustering. You'll learn how to:

1. Preprocess data for clustering
2. Perform and visualize dimensionality reduction
3. Implement K-means clustering from scratch
4. Evaluate clustering results

Follow the instructions in each section and fill in the code where indicated.

```
# Load Data and Preprocess
def load_data(filepath):
    pass
#===== FOR PCA Dimensionality reduction =====
# Apply PCA for Dimensionality Reduction
def apply_pca(x, n_components):
    pass

#===== FOR K-Means =====
# Find Optimal Clusters for KMeans (Elbow Method)
def find_optimal_clusters(x, max_clusters=10):
    pass

# Perform KMeans Clustering
# Change None to the number of n_clusters value from the elbow method
def perform_kmeans_clustering(x, n_clusters=None):
    pass

#===== FOR Agglomerative =====
# Perform Agglomerative Clustering
# Change None to the number of n_clusters value from the elbow method
def perform_agglomerative_clustering(x, n_clusters=None):
    pass

#=====FOR Dendrogram=====
# Get Linkages for Dendrogram
def get_linkages(x):
    pass

# Plot Dendrogram
def plot_dendrogram(linked):
    pass
```

▼ 1. Data Loading and Preprocessing

First, complete the data preprocessing function below. You need to:

1. Load the data
2. Handle categorical variables
3. Scale numerical features

```
def load_data(filepath):
    """Load and preprocess the bank marketing dataset.

    TODO:
        1. Load the CSV file (hint: it uses semicolon separator)
```



```

2. Convert categorical columns to numerical using LabelEncoder
3. Scale the features using StandardScaler
"""

# Your code here:
# Load data
# df = # TODO: Load CSV with correct separator
df = pd.read_csv(filepath, sep=';')

# List of categorical columns to encode
categorical_cols = ['job', 'marital', 'education', 'default', 'housing',
                    'loan', 'contact', 'month', 'poutcome', 'y']

# TODO: Apply label encoding to categorical columns
le = LabelEncoder()
for col in categorical_cols:
    df[col] = le.fit_transform(df[col])

# Select features for clustering
features = ['age', 'balance', 'campaign', 'previous', 'job', 'education',
            'housing', 'loan', 'default']
# X = # TODO: Select features
X = df[features]

# TODO: Scale the features
# X_scaled = # TODO: Apply StandardScaler
scalar = StandardScaler()
X_scaled = scalar.fit_transform(X)

return X_scaled, df

```

```

# Load and preprocess the data
X_scaled, data = load_data('bank-full.csv')
print("Data shape:", X_scaled.shape)

```

```
Data shape: (45211, 9)
```

```

class KMeansClustering:
    """Minimal K-means skeleton for students to implement.

    Students should implement the core methods below:
    - _initialize_centroids
    - _assign_clusters
    - _update_centroids
    - fit

    Keep implementations simple and readable; tests and visualization code will
    use these methods once implemented.
    """

    def __init__(self, n_clusters=3, max_iters=100, random_state=42):
        self.n_clusters = n_clusters
        self.max_iters = max_iters
        self.random_state = random_state
        self.centroids = None
        self.labels = None

    def _initialize_centroids(self, X):
        """Initialize centroids.

        TODO (student):
        - Randomly select `n_clusters` distinct points from X as initial centroids.
        - Return an array of shape (n_clusters, n_features).
        Hint: Use np.random.choice to pick indices.
        """

        np.random.seed(self.random_state)
        n_samples, n_features = X.shape
        indices = np.random.choice(n_samples, self.n_clusters, replace=False)
        centroids = X[indices]
        return centroids
        # raise NotImplementedError("_initialize_centroids must be implemented by the student")

    def _assign_clusters(self, X):
        """Assign each sample in X to the nearest centroid.

        TODO (student):
        - Compute distance from each point to each centroid (Euclidean)

```



```

- Return an integer array of shape (n_samples,) with cluster labels
Hint: np.linalg.norm with axis manipulation or broadcasting helps here.
"""

n_samples = X.shape[0]
distances = np.zeros((n_samples, self.n_clusters))

for i in range(self.n_clusters):
    distances[:, i] = np.linalg.norm(X - self.centroids[i], axis=1)

labels = np.argmin(distances, axis=1)
return labels
# raise NotImplementedError("_assign_clusters must be implemented by the student")

def _update_centroids(self, X, labels):
    """Recompute centroids as the mean of points assigned to each cluster.

    TODO (student):
    - For each cluster id in 0..n_clusters-1 compute the mean of points
      assigned to that cluster. If a cluster has no points, consider reinitializing
      its centroid (or leave unchanged) – discuss in your report.
    - Return an array of shape (n_clusters, n_features).
    """
    n_features = X.shape[1]
    new_centroids = np.zeros((self.n_clusters, n_features))

    for i in range(self.n_clusters):
        cluster_points = X[labels == i]

        if len(cluster_points) > 0:
            new_centroids[i] = cluster_points.mean(axis=0)
        else:
            # Handle empty cluster: re-initialize
            indices = np.random.choice(X.shape[0], 1, replace=False)
            new_centroids[i] = X[indices]

    return new_centroids
# raise NotImplementedError("_update_centroids must be implemented by the student")

def fit(self, X):
    """Run K-means until convergence or max_iters.

    TODO (student):
    - Initialize centroids
    - Loop: assign clusters, update centroids
    - Stop early if centroids do not change (or change below a tiny threshold)
    - Store final labels in self.labels and centroids in self.centroids
    - Return self
    """
    self.centroids = self._initialize_centroids(X)

    for _ in range(self.max_iters):
        self.labels = self._assign_clusters(X)
        new_centroids = self._update_centroids(X, self.labels)

        if np.allclose(self.centroids, new_centroids, atol=1e-4):
            break

        self.centroids = new_centroids

    self.labels = self._assign_clusters(X) # Final assignment
    return self
# raise NotImplementedError("fit must be implemented by the student")

def predict(self, X):
    """Assign cluster labels to X using the learned centroids.

    Implementation may call _assign_clusters but should error if centroids
    are not yet initialized (i.e., if fit wasn't called).
    """
    if self.centroids is None:
        raise ValueError("Model has not been fitted yet. Call fit(X) first.")
    return self._assign_clusters(X)

```

2. Dimensionality Reduction



Before clustering, we often reduce the dimensionality of our data for better visualization and performance. Implement PCA below:

```
def apply_pca(X, n_components=2):
    """Apply PCA for dimensionality reduction.

    TODO:
    1. Initialize and fit PCA
    2. Transform the data
    3. Create visualizations to understand:
        - Explained variance ratio
        - Cumulative explained variance
        - Data distribution in 2D
    """

    # Your code here:
    pca = PCA(n_components=n_components)
    X_pca = pca.fit_transform(X)

    # TODO: Create visualization
    plt.figure(figsize=(12, 4))

    # Plot explained variance
    plt.subplot(1, 3, 1)
    plt.bar(range(1, len(pca.explained_variance_ratio_)+1), pca.explained_variance_ratio_)
    plt.title('Explained Variance Ratio')
    plt.xlabel('Principal Component')
    plt.ylabel('Variance Ratio')


    # Plot cumulative explained variance
    plt.subplot(1, 3, 2)
    plt.plot(range(1, len(pca.explained_variance_ratio_)+1), np.cumsum(pca.explained_variance_ratio_))
    plt.title('Cumulative Explained Variance')
    plt.xlabel('Number of Components')
    plt.ylabel('Cumulative Variance')

    # Plot data in 2D
    plt.subplot(1, 3, 3)
    plt.scatter(X_pca[:, 0], X_pca[:, 1], alpha=0.5)
    plt.title('Data in 2D PCA Space')
    plt.xlabel('PC1')
    plt.ylabel('PC2')

    plt.tight_layout()
    plt.show()

    return X_pca

# Apply PCA
X_pca = apply_pca(X_scaled)
print("Shape after PCA:", X_pca.shape)
```



Shape after PCA: (45211, 2)

3. Clustering Evaluation

Implement functions to evaluate the quality of your clustering results:

```

def calculate_inertia(X, labels, centroids):
    """Calculate the within-cluster sum of squares (inertia).

    TODO:
    1. For each cluster, calculate the sum of squared distances
       between points and their centroid
    2. Sum up all cluster distances
    """
    # Your code here
    inertia = 0
    for k in range(len(centroids)):
        points = X[labels == k]
        if len(points) > 0:
            inertia += np.sum((points - centroids[k])**2)
    return inertia

def plot_elbow_curve(X, max_k=10):
    """Plot the elbow curve to find optimal number of clusters.

    TODO:
    1. Try different values of k (1 to max_k)
    2. Calculate inertia for each k
    3. Plot k vs inertia
    4. Help identify the 'elbow' point
    """
    # Your code here
    inertias = []
    K = range(1, max_k + 1)

    # Calculate inertia for different k values
    for k in K:
        kmeans = KMeansClustering(n_clusters=k, random_state=42)
        kmeans.fit(X)
        inertia = calculate_inertia(X, kmeans.labels, kmeans.centroids)
        inertias.append(inertia)

    # Create elbow plot
    plt.figure(figsize=(8, 5))
    plt.plot(K, inertias, 'bx-')
    plt.xlabel('Number of clusters (k)')
    plt.ylabel('Inertia')
    plt.title('Elbow Method for Optimal k')
    plt.grid(True)
    plt.show()

    return inertias

# Try different numbers of clusters
inertias = plot_elbow_curve(X_pca)

# Apply final clustering
kmeans = KMeansClustering(n_clusters=3) # Try different values based on elbow curve
kmeans.fit(X_pca)

# Visualize final results
plt.figure(figsize=(12, 4))

# Plot clusters
plt.subplot(1, 2, 1)
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=kmeans.labels, cmap='viridis')
plt.scatter(kmeans.centroids[:, 0], kmeans.centroids[:, 1],
           c='red', marker='x', s=200, label='Centroids')
plt.title('Final Clustering Results')
plt.legend()

# Plot evaluation metrics
plt.subplot(1, 2, 2)
# Your code here: Add relevant evaluation metric plots
plt.text(0.1, 0.5, f'Inertia: {calculate_inertia(X_pca, kmeans.labels, kmeans.centroids):.2f}\nSilhouette: {silhouette_score(X_pca, kmeans.labels):.2f}')
plt.title('Evaluation Metrics')
plt.axis('off')


plt.tight_layout()
plt.show()

```



```
# Calculate and print evaluation metrics
inertia = calculate_inertia(X_pca, kmeans.labels, kmeans.centroids)
silhouette = silhouette_score(X_pca, kmeans.labels)

print("\nClustering Evaluation:")
print(f"Inertia: {inertia:.2f}")
print(f"Silhouette Score: {silhouette:.2f}")
```



4. Recursive Bisecting K-means (Student exercise)

This is an optional exercise for students who want to explore a hierarchical variant of K-means.

Task: implement a concise bisecting K-means procedure that recursively splits clusters into two until a target number of clusters is reached.

Learning goals:

- Understand how repeated binary splits can form a hierarchical clustering
- Practice applying K-means on subclusters and tracking labels/centroids

Hints:

- You can use sklearn's KMeans(k=2) for the binary split step, or reuse your `KMeansClustering` implementation.
- Keep label bookkeeping simple: use increasing integer labels for new clusters.



- Store split metadata (parent \rightarrow left/right) to enable a tree visualization later.

```

class BisectingKMeans:
    """Concise skeleton for students to implement a bisecting K-means algorithm.

    Students should implement `fit_predict` to recursively split clusters until
    `n_clusters` is reached.
    """
    def __init__(self, n_clusters, random_state=42):
        self.n_clusters = n_clusters
        self.random_state = random_state
        self.labels_ = None
        # Optional: store mapping parent  $\rightarrow$  (left_label, right_label)
        self.split_tree = {}
        # Optional: store centroids per cluster id
        self.centers_ = {}

    def fit_predict(self, X):
        """Recursively bisect clusters until `n_clusters` is reached.

        TODO (student):
        - Start with all points assigned to label 0.
        - While number of unique labels < n_clusters:
            - Select a cluster to split (e.g., the largest cluster by size)
            - Run a binary KMeans ( $k=2$ ) on the points in that cluster
            - Assign new labels (keep one child label as the original, give the other a new id)
            - Record parent  $\rightarrow$  (left, right) in `self.split_tree` and centroids in `self.centers_`
        - Set and return `self.labels_` (numpy array of length n_samples)

        Hints:
        - Use sklearn.cluster.KMeans(n_clusters=2, random_state=self.random_state) for the split step
        - Keep a counter for new label ids and increment when creating a new cluster
        - Use boolean indexing to operate on subsets of X efficiently
        """
        n_samples = X.shape[0]
        # 1. Start with all points assigned to label 0
        self.labels_ = np.zeros(n_samples, dtype=int)

        active_clusters = [0]
        next_label_id = 1
        self.centers_[0] = X.mean(axis=0)

        # 2. Loop until we have the desired number of clusters
        while len(active_clusters) < self.n_clusters:

            # 3. Select the largest cluster to split
            largest_cluster_label = -1
            max_size = -1
            for label in active_clusters:
                size = np.sum(self.labels_ == label)
                if size > max_size:
                    max_size = size
                    largest_cluster_label = label

            if max_size < 2:
                break # Stop if clusters are too small to split

            cluster_mask = (self.labels_ == largest_cluster_label)
            cluster_points = X[cluster_mask]

            # 4. Run binary KMeans on these points
            kmeans_split = SklearnKMeans(n_clusters=2, random_state=self.random_state, n_init=10)
            kmeans_split.fit(cluster_points)

            # 5. Assign new labels
            original_indices = np.where(cluster_mask)[0]
            split_labels = kmeans_split.labels_

            child_label_1 = largest_cluster_label
            child_label_2 = next_label_id
            next_label_id += 1

            self.labels_[original_indices[split_labels == 0]] = child_label_1
            self.labels_[original_indices[split_labels == 1]] = child_label_2

            # 6. Record split metadata
            self.split_tree[largest_cluster_label] = (child_label_1, child_label_2)

```



```
self.split_tree[largest_cluster_label] = (child_label_1, child_label_2)
self.centers_[child_label_1] = kmeans_split.cluster_centers_[0]
self.centers_[child_label_2] = kmeans_split.cluster_centers_[1]

active_clusters.remove(largest_cluster_label)
active_clusters.append(child_label_1)
active_clusters.append(child_label_2)

return self.labels_
# raise NotImplementedError('Implement the bisecting algorithm as a student exercise')

# Example (for instructor use only):
# bisect = BisectingKMeans(n_clusters=4)
# labels = bisect.fit_predict(X_pca)
# plt.scatter(X_pca[:,0], X_pca[:,1], c=labels)
```



```
from sklearn.cluster import KMeans as SklearnKMeans
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import numpy as np

print("Running Bisecting K-Means...")

bisect = BisectingKMeans(n_clusters=3, random_state=42)

labels = bisect.fit_predict(X_pca)

final_labels = np.unique(labels)
final_centroids_list = [bisect.centers_[label] for label in final_labels]
final_centroids_array = np.array(final_centroids_list)


plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:,0], X_pca[:,1], c=labels, cmap='viridis', s=10, alpha=0.7)

plt.scatter(final_centroids_array[:, 0], final_centroids_array[:, 1],
           c='red', marker='X', s=200, label='Centroids', edgecolor='black')

plt.title('Recursive Bisecting K-means Results (k=3)')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.legend()
plt.savefig("bisecting_kmeans_clusters.png")
plt.show()

silhouette_bisect = silhouette_score(X_pca, labels)

print(f"\nFinal cluster labels created: {np.unique(labels)}")
print(f"Silhouette Score (Bisecting K-Means): {silhouette_bisect:.2f}")
```



Remember to document your code and explain your findings!